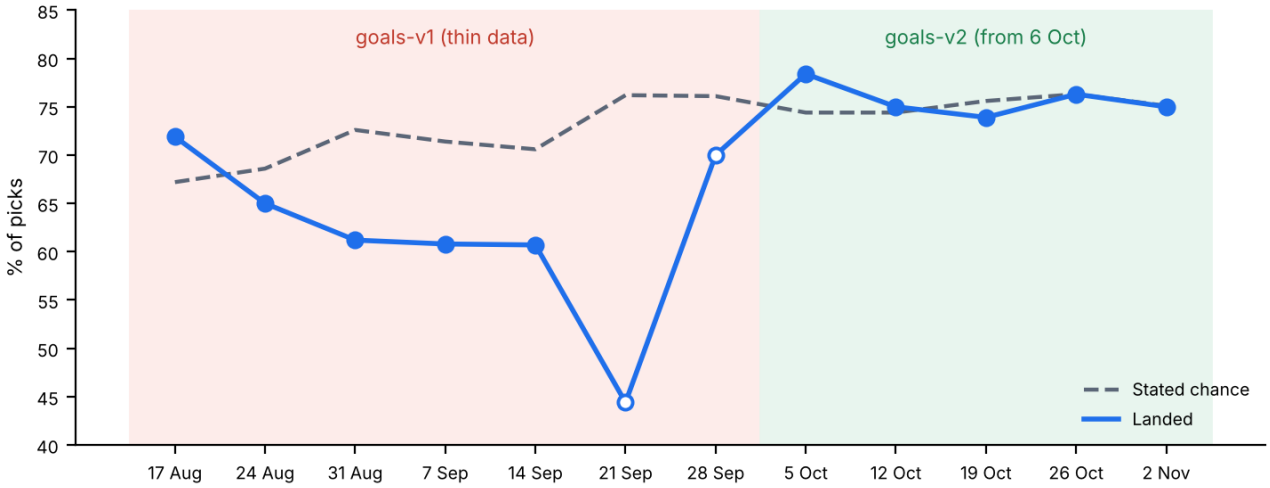
Model audit: what we fixed and what it did

Team markets only (no player props). Prepared 4 November 2026. Code is on branch feat/model-performance-training.

Bottom line. The decline you saw was mostly the old model (goals-v1) working from thin data. It was already fixed on 6 Oct when goals-v2 took over: 75.8% landed against 75.2% stated on 1,032 picks since. What was left is smaller, and we acted on it: from **5 Nov** the model drops its weakest picks (under 65%), leans toward overs as you asked, and holds back the CAF Champions League. Across 17,857 test games that lifts the hit rate from 77.2% to 77.8%, and on your September and October games from 78.0% to 78.4%. **Today's picks are not touched.**

1. Why picks declined



Weekly live record. Hollow dots are weeks with fewer than 20 settled picks. Shaded: red is the old model, green the new one.

Live record (settled picks)	Landed	Stated	Picks
goals-v1, leagues in the catalogue (to 6 Oct)	66.3%	69.9%	267
goals-v1, no league code (trained on 15-35 results)	57.8%	72.3%	166
goals-v2 (since 6 Oct)	75.8%	75.2%	1,032
Under picks, Sep-Oct (paused 16 Oct)	56.1%	72.4%	107

Replaying goals-v2 on the September and October fixtures gives 76.8% and 78.9%, so September would not have dipped under the current model.

2. How it was tested

Every number comes from a walk-forward test: each game is predicted using only results from before its matchday, through the same code the site runs, on 17 European divisions. Two samples, as you asked:

	Sample	Used for
A (mine)	Aug 2023 to Jun 2026, 17,857 games. Split: rules derived on 2023-25, scored on the 2025-26 season (5,911 games).	Choosing between options. Every decision was made here.
B (yours)	All of September and October 2026, 1,165 games.	Final check only. Nothing was tuned on it.

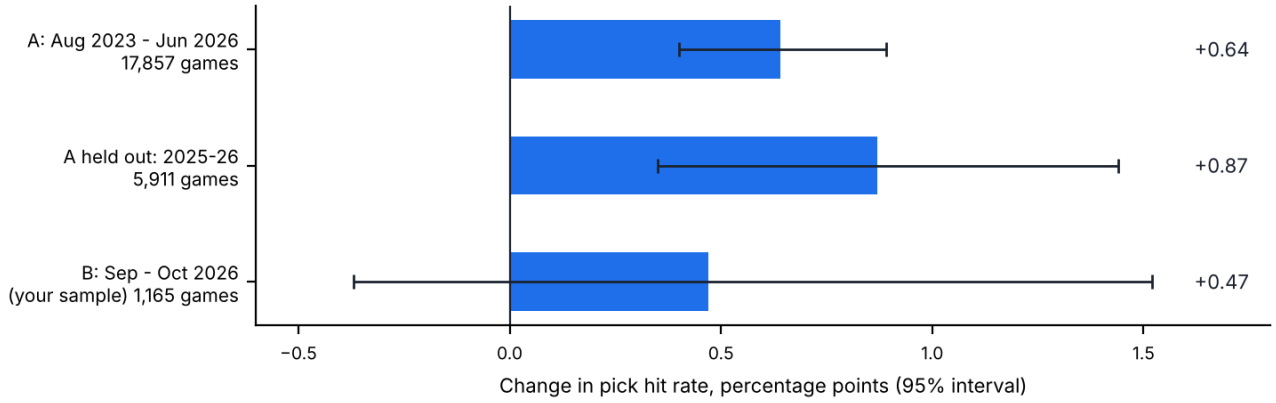
Caveat: B was seen once in aggregate during the earlier audit sweep, before it was set aside as a check sample. Sample A includes 2024-25, the season the old settings were tuned on, which favours the old model if anything.

3. What ships on 5 Nov

Change	Why
No headline pick under 65%	Ranker picks stated 60-65% landed 60.5% (A held out, n=243), 61.0% (B, n=41) and 63.3% (live since 7 Oct, n=49). Picks at 65% and up landed at or above stated. About 4% of games lose their pick.
Overs lean (+0.02 in the ranking)	Your call. On the replay it is neutral: +0.16 (A) and 0.00 (B) points, intervals include zero. It adds roughly 30% more over picks. A stronger lean (0.04) and a lower bar for over 2.5 were tested and were not better.
CAF Champions League: no headline pick	9 of 23 live picks landed against 64% stated; history is thin and pooled. Reversible in one line (params.ts). No replay data exists for this competition.
Neutral-venue totals	Tournament totals were pulled toward a mean that includes home advantage, about 2.8% too high. Fixed and unit tested.
Promoted-club prior off for national-team and cup pools	It was only ever tested on domestic leagues.
Finished games show pre-match numbers	Match pages used to refit with the game's own result in the data. Fixed and unit tested.

Today's picks are excluded. Every game kicking off before 00:00 Lagos on 5 Nov uses exactly the old parameters. The gate is on kickoff time, not on when a prediction is computed, because pages and the logging job rebuild predictions on every run. A gate on 'now' would have rewritten today's picks the next time they were rebuilt. Tests pin the boundary to the millisecond (22:59:59Z is old, 23:00:00Z is new) and check that the old settings predict identically. Tomorrow's slate is the first to use the new settings, so tomorrow's picks, if already visible, may change.

4. Results



Sample	Hit before	Hit after	Change (95% interval)	Picks kept	Overs in picks
A, 17,857 games	77.2%	77.8%	+0.64 (0.40 to 0.89)	96.5%	23% to 29%
A held out, 2025-26	77.1%	78.0%	+0.87 (0.35 to 1.44)	96.3%	24% to 30%
B, Sep-Oct 2026	78.0%	78.4%	+0.47 (-0.37 to 1.52)	97.0%	32% to 37%

Read these honestly: the gain is modest, about half to nine-tenths of a point. It is solid on the 17,857-game sample and its 2025-26 slice, and points the same way on your 1,165 games but is too small a sample to be sure on its own. The probabilities themselves are unchanged by this release; the gain comes from which picks go out. The over picks that were added land slightly less often (77.4% against 77.8%) than the over picks before, but better than the picks they replaced.

5. Tried and not shipped

Change	What happened
Lower regularisation plus a wider gap between strong and weak sides	Probabilities got better (1X2 log loss down 0.0024 to 0.0028, interval excludes zero) but picks got worse: -0.7 to -1.2 points. Refitting the pick rules on its own numbers and scoring the next season gave 76.4% against 77.3% for the current model refitted the same way. Better probabilities, not better picks.
Adjusting for recent scoring level (60 or 90-day)	No gain on sample A (over/under log loss 0.6821 against 0.6820); a small gain only on Sep-Oct.
Lowering the over 2.5 bar from 70% to 65%	Hit rate -0.69 points on B.

6. Correction to the earlier audit, and what is still open

Correction. The earlier audit called quarter-line grading a bug. It is not: an existing test shows a half win counting as a win is intended, and plain-line totals agree. Only a code comment was wrong. I had changed the grading, the test caught it, and I reverted it.

Still open	Detail
Extra-time scores in national-team data	The source records 120-minute scores and has no flag to separate them. Needs another source.
Corners picks	75.0% landed against 81.8% stated live (n=48). Watch it; not enough to act on.
Liga MX	62.8% against 74.0% stated (n=43). Watch.
Market check evidence	Backtested on closing prices, which are sharper than the price at publish. Needs a few hundred publish-time prices before it can be trusted.
Frozen holdout	From now on, nothing should be tuned on games after 1 Nov 2026, so the next change has a clean test.

Reproduce: scripts/audit-eval.ts, audit-report.ts and audit-layer.ts (see docs/model-audit-teams-2026-11.md, section G). All 1,220 tests pass.